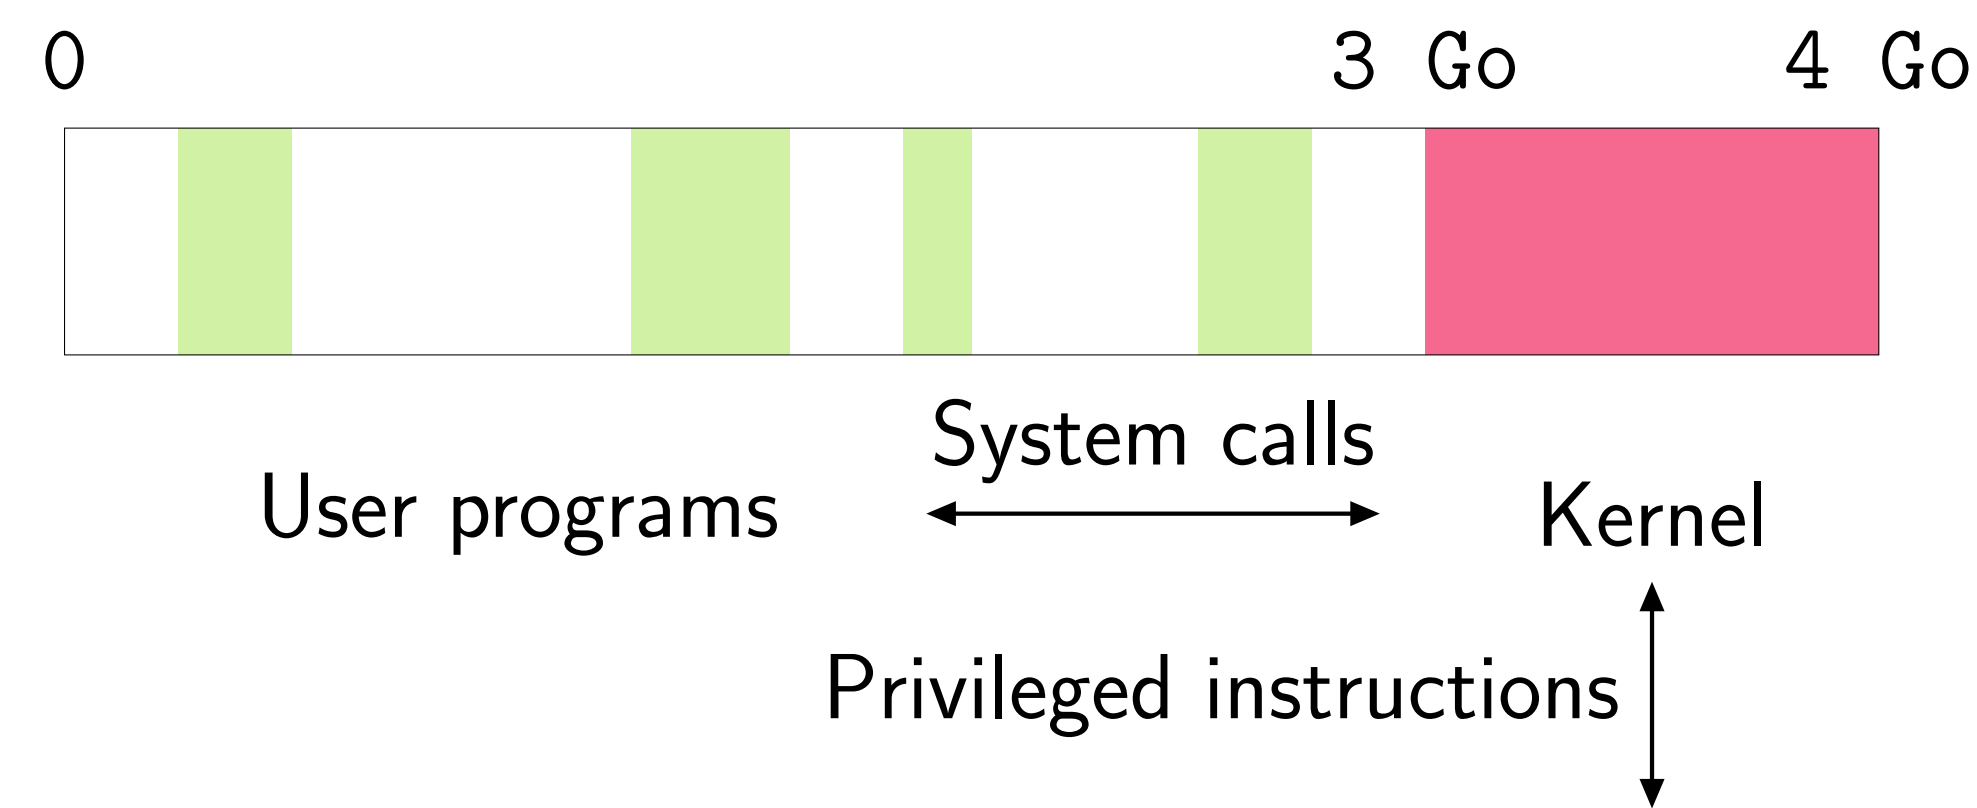


Verifying the Safety of User Pointers Using Static Typing


Etienne Millon^{1,2} Emmanuel Chailloux¹ Sarah Zennou²

¹Sorbonne Universités, UPMC Univ Paris 06, UMR 7606, LIP6, F-75005, Paris, France ² Airbus Group Innovations

Separation in operating systems



Only the kernel can access the hardware.
User programs have to use *system calls* (open, read, ...).



Execution of a piece of code:

- ▶ with processor at level P
 - ▶ with privilege level C
 - ▶ accessing data with level D
- is possible iff

$$P \leq \min\{C; D\}$$

On x86, levels are *rings*:
Kernel = 0 and User = 3.

P	C	D	Access
Kernel	Kernel	Kernel	OK
Kernel	Kernel	User	OK
Kernel	User	Kernel	OK
Kernel	User	User	OK
User	Kernel	Kernel	–
User	Kernel	User	–
User	User	Kernel	–
User	User	User	OK

The confused deputy problem

User programs can pass structures to system calls by pointer:

```
struct timeval tv;
int z = gettimeofday(&tv, NULL);
```

The kernel fills in tv with its own privileges (not the caller's).

A user can write the current time of the day at any address.

This is *the confused deputy problem*.

The kernel should check at run time that pointers controlled by userspace point to userspace (copy_{from,to}_user).

We detect the places where this dynamic check is omitted, using static typing of C code.

Pointer types

Depending on who controls their value, i.e. how they are created:

- ▶ kernel pointers: &x, etc. They can be dereferenced.

$$\frac{\Gamma \vdash e : t *}{\Gamma \vdash *e : t}$$

- ▶ user pointers: system call arguments. They need a dynamic check.

$$\frac{\Gamma \vdash e_1 : t * \quad \Gamma \vdash e_2 : t @}{\Gamma \vdash \text{copy from user}(e_1, e_2) : \text{INT}}$$

User pointer sources require annotations: one per system call.

Example: freedesktop.org bug #29340

```
int radeon_info_ioctl(struct drm_device *dev,
                     void *data,
                     struct drm_file *filp) {
    /*!nprk userptr_fieldp data value*/
    struct drm_radeon_info *info = data;
    uint32_t *value_ptr = (uint32_t *)
        ((unsigned long)info->value);
    uint32_t value = *value_ptr;
    /* ... */
}
```

Inference output – **error**:

```
05-drm.c:17#10 - Type clash between :
  KPTr (_a15)
  UPTr (_a8)
```

But if we replace last line by the following:

```
if (copy_from_user(&value, value_ptr,
                  sizeof(value)))
    return -14;
```

Inference output – **fully annotated program**:

```
(06-drm-ok.c:17#6) ^{
  Int tmp_cir!0;
  (06-drm-ok.c:17#6) ^tmp_cir!0 <-
    copy_from_user
      ( &(value) : KPTr (d),
        value_ptr_UPTr (d) : UPTr (d),
        4 : Int); }
```

Compilation

```
int f (int* x) {
  return (*x + 1);
}
```

Type erasure

```
let f =
  fun (x) ->
    return (*x + 1)
```

Type inference

```
let f : Int* -> Int =
  fun (x : Int) ->
    return ((((*x) : Int)
      + (1 : Int))) : Int)
```

Input is C code.
GNU extensions used in the kernel are handled.

C type annotations are removed.
The intermediate language has first-order functions and left-values (for partial updates).

Every subexpression gets a type.
This is more precise than C types since abstract types (e.g. user pointers) can be inferred.

Bibliography

■ N. Hardy.

The confused deputy (or why capabilities might have been invented).
ACM Operating Systems Review, 1988.

■ C. Hymans and O. Levillain.

Newspeak, Doubleplussimple Minilang for Goodthinkful Static Analysis of C.
Technical Note 2008-IW-SE-00010-1, EADS IW/SE, 2008.

■ R. Johnson and D. Wagner.

Finding user/kernel pointer bugs with type inference.
In *USENIX Security Symposium*, 2004.